

Name- Harshita Kiran Chaudhari

Contact No- 9167712338

Email : harshitachaudhari552006@gmail.com

Eiger AI Security Assessment

Hands-on AI Security Assessment Report

Assessment Type: AI Security / Adversarial Testing

Environment: Eiger AI Security Lab

Assessment Coverage: M1 – Prompt Injection + M2 – Output Handling

Testing Type: Authorized security assessment in a local training environment

1. Executive Summary

This document presents my hands-on security assessment of the Eiger AI Security Lab.

The assessment focused on identifying, reproducing, and validating security weaknesses in two Eiger exercises:

- **M1 – Prompt Injection**
- **M2 – Output Handling / Stored XSS**

For each exercise, I analyzed the vulnerable implementation, reproduced the security issue in the authorized local lab environment, observed the resulting impact, applied the available security control, and validated the behavior before and after remediation.

The assessment demonstrated two important security boundaries:

- AI systems must not expose secrets through attacker-controlled prompts or model-accessible instructions.
- Web applications must not render user-controlled data as trusted HTML.

The results were verified using the Eiger validation endpoints.

2. Assessment Methodology

The following methodology was followed:

1. Set up the Eiger environment locally using Docker.
2. Reviewed the relevant application code and security controls.
3. Identified the vulnerability and its security impact.
4. Reproduced the vulnerability in the authorized lab.
5. Captured the vulnerable behavior and validation result.
6. Applied the relevant security control.

7. Repeated the same attack after remediation.
 8. Compared the before and after results.
 9. Verified the remediation using the Eiger validation mechanism.
 10. Documented the evidence and AI assistance used.
-

3. M1 – Prompt Injection

Objective

Demonstrate whether attacker-controlled instructions could manipulate the AI system into disclosing sensitive internal information.

Vulnerability / Security Concept

Prompt Injection → Sensitive Information Disclosure

The vulnerable implementation combined the system instructions, internal token, and user-controlled message into a model-accessible prompt.

The internal operator token used in the lab was:

HALCYON-OPS-7731

Because the secret was present in the model-accessible prompt context, an attacker could use prompt injection to influence the model into disclosing it.

Steps to Reproduce

Step 1 – Open M1

Open the M1 exercise in the Eiger environment.

Step 2 – Interact with Halo

Interact with the vulnerable AI chatbot.

Step 3 – Perform Prompt Injection

Use a prompt-injection request attempting to reveal the internal instructions and secret token.

Step 4 – Observe the Response

The vulnerable model disclosed the internal operator token.

Step 5 – Validate

The Eiger validation endpoint confirmed the successful attack.

Observed Result – Before Remediation

The prompt injection successfully caused the internal operator token to be disclosed.

Validation:

Core: PASS

Stretch: FAIL

Remediation

The system prompt hardening control was enabled:

SEC_SYSTEM_PROMPT_HARDENING=on

The hardened implementation separates system instructions from the user message instead of combining them into one user-controlled prompt.

The secret token is also removed from the model-accessible prompt when the security control is enabled.

Observed Result – After Remediation

The same attack was repeated after remediation.

The internal operator token was no longer disclosed.

Validation:

Core: FAIL

Stretch: FAIL

Impact

Successful prompt injection can potentially expose:

- System instructions
- Internal tokens
- Credentials
- Confidential application information
- Sensitive contextual information

The risk becomes higher when the AI system has access to privileged tools or sensitive enterprise data.


Prevention / Mitigation

Recommended controls include:

- Treat user input as untrusted.
 - Separate system instructions from user input.
 - Never place secrets inside model-accessible prompts.
 - Apply input and output validation.
 - Use least-privilege access controls.
 - Perform adversarial prompt-injection testing.
-

Evidence

M1 Before – Successful Prompt Injection



```
Windows PowerShell
Container eiger-mcp-core-banking-1 Running 0.0s
Container eiger-web-1 Started 8.2s
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker compose ps
NAME                IMAGE                                COMMAND                                SERVICE    CREATE
eiger-db-1          postgres:16                         "docker-entrypoint.s..."            db         2 hour
s ago              Up 42 minutes (healthy)            5432/tcp
eiger-mcp-core-banking-1 eiger-mcp-core-banking            "uv run uvicorn halc..."            mcp-core-banking 2 hour
s ago              Up 42 minutes                    127.0.0.1:9001->9001/tcp
eiger-mcp-crm-1     eiger-mcp-crm                      "uv run uvicorn halc..."            mcp-crm        2 hour
s ago              Up 42 minutes                    127.0.0.1:9002->9002/tcp
eiger-ollama-1     ollama/ollama:latest              "/bin/ollama serve"                  ollama         2 hour
s ago              Up 42 minutes                    11434/tcp
eiger-web-1         eiger-web                          "uv run uvicorn halc..."            web            About
a minute ago      Up About a minute                127.0.0.1:8000->8000/tcp
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker exec eiger-web-1 env | findstr /i "SEC_SYSTEM_PROMPT
_HARDENING"
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker exec eiger-web-1 env | findstr /i "SEC_SYSTEM_PROMPT
_HARDENING"
PS C:\Users\Harshu\OneDrive\Desktop\eiger> Invoke-WebRequest "http://127.0.0.1:8000/validate/m1?sessio
n=eiger-0e97c839" -UseBasicParsing | Select-Object -ExpandProperty Content
{"core":"pass","stretch":"fail"}
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker exec eiger-web-1 env | findstr /i "SEC_SYSTEM_PROMPT
_HARDENING"
PS C:\Users\Harshu\OneDrive\Desktop\eiger>
```

M1 After – Remediated

```
Windows PowerShell
OL_CALL)
halcyon\validators\m5.py:10: progress.mark(store, session_id, MODULE, core, stretch)
halcyon\validators\m5.py:12: "core": "pass" if core else "fail",
halcyon\validators\m6.py:8: core = audit.has_event(store, session_id, MODULE, audit.MCP_POISONED_IN
VOCATION)
halcyon\validators\m6.py:11: progress.mark(store, session_id, MODULE, core, stretch)
halcyon\validators\m6.py:13: "core": "pass" if core else "fail",
halcyon\validators\m7.py:8: core = (audit.has_event(store, session_id, MODULE, audit.INTER_AGENT_IN
JECTION_PROPAGATED)
halcyon\validators\m7.py:11: progress.mark(store, session_id, MODULE, core, stretch)
halcyon\validators\m7.py:13: "core": "pass" if core else "fail",
halcyon\validators\m8.py:8: core = audit.has_event(store, session_id, MODULE, audit.GUARDRAIL_BYPAS
SED)
halcyon\validators\m8.py:10: progress.mark(store, session_id, MODULE, core, stretch)
halcyon\validators\m8.py:12: "core": "pass" if core else "fail",
halcyon\validators\m8.py:14: _mcp_core_url = os.environ.get("MCP_CORE_URL"),
halcyon\validators\m8.py:16: def _mcp_up(core_url: str, crm_url: str) -> bool:
halcyon\validators\m8.py:18:     for url in (core_url, crm_url):
halcyon\validators\m8.py:20:         if not (_mcp_core_url and _mcp_crm_url):
halcyon\validators\m8.py:22:             return "up" if _mcp_up(_mcp_core_url, _mcp_crm_url) else "down"
halcyon\validators\m8.py:24:         "core": result.get("core", "fail"),
PS C:\Users\Harshu\OneDrive\Desktop\eiger> curl.exe "http://127.0.0.1:8000/validate/m1?session=NEW_SE
SION_ID"
{"core":"fail","stretch":"fail"}
PS C:\Users\Harshu\OneDrive\Desktop\eiger>
```

4. M2 – Output Handling

Objective

Demonstrate how unsafe handling of a stored display name can allow attacker-controlled HTML/JavaScript content to execute when the value is rendered by the browser.

Vulnerability / Security Concept

Improper Output Encoding / Stored Cross-Site Scripting (XSS)

The vulnerable application accepts a user-controlled profile display name and renders it in the web page.

When output encoding is disabled, the value is returned unchanged:

```
def encode_output(text: str, settings: Settings) -> str:
    if settings.sec_output_encoding:
        return html.escape(text)
    return text
```

Therefore, attacker-controlled HTML can reach the browser as active content.

Steps to Reproduce

Step 1 – Open M2

Open the M2 Output Handling exercise in vulnerable mode.

Step 2 – Identify the Input

The affected input is the stored profile display name.

Step 3 – Supply XSS Payload

A test payload was used:

```
<img src=x onerror=fetch('/beacon/xss?session=<SESSION>')>
```

Step 4 – Save the Display Name

The payload was submitted as the profile display name using the profile API.

Step 5 – Open the Chat Page

The page was opened with the corresponding Eiger session.

Step 6 – Observe Execution

Because the value was not encoded, the browser interpreted the supplied value as HTML.

The onerror handler triggered the Eiger XSS beacon.

Step 7 – Validate

The Eiger validator recorded the XSS beacon.

Observed Result – Before Remediation

The XSS payload successfully executed.

Validation:

Core: PASS

Stretch: FAIL

Remediation

Output encoding was enabled:

SEC_OUTPUT_ENCODING: "on"

This activates:

```
return html.escape(text)
```

The same XSS payload was then tested again.

Observed Result – After Remediation

The same attack was repeated after enabling output encoding.

The injected HTML was no longer interpreted as executable HTML.

Validation:

Core: FAIL

Stretch: FAIL

Why the Fix Works

html.escape() converts HTML special characters into safe HTML entities.

For example:

```
<img src=x onerror=...>
```

is converted so that the browser treats it as **text rather than an HTML element**.

Therefore, the injected event handler does not execute.

Impact

Successful stored XSS can potentially allow attacker-controlled browser-side code to execute in the victim's browser.

Potential impacts include:


- Session or sensitive-data exposure
- Unauthorized browser actions
- Page-content manipulation
- Phishing
- Access to browser-accessible application information
- User-interface manipulation

Prevention / Mitigation

Recommended controls include:

- Apply contextual output encoding.
- Never treat user-controlled values as trusted HTML.
- Avoid unnecessary safe HTML rendering.
- Use automatic framework escaping.
- Implement a restrictive Content Security Policy.
- Validate security-sensitive stored fields.
- Perform security testing on user-controlled inputs.
- Use defense in depth.

M2 Before – Stored XSS



```
cmn common_init.: warming up the model with an empty run - please wait ... (--no-warmup to disable)
time=2026-10-07T12:15:48.134Z level=INFO source=llama_server.go:1362 msg="waiting for llama-server to
become available" status="llm server not responding"
time=2026-10-07T12:15:48.754Z level=INFO source=llama_server.go:1362 msg="waiting for llama-server to
become available" status="llm server loading model"
time=2026-10-07T12:15:49.233Z level=INFO source=llama_server.go:1362 msg="waiting for llama-server to
become available" status="llm server not responding"
time=2026-10-07T12:15:49.522Z level=INFO source=llama_server.go:1362 msg="waiting for llama-server to
become available" status="llm server loading model"
time=2026-10-07T12:15:52.532Z level=WARN source=llama_server.go:1313 msg="client connection closed bef
ore llama-server finished loading, aborting load"
time=2026-10-07T12:15:52.541Z level=INFO source=sched.go:646 msg="Load failed" model=/root/.ollama/mod
els/blobs/sha256-dde5aa3fc5ffc17176b5e8bdc82f587b24b2678c6c66101bf7da77af9f7ccdf error="timed out wai
ting for llama-server to start: context canceled"
[GIN] 2026/10/07 - 12:15:53 | 499 | 2m8s | 172.18.0.6 | POST "/api/generate"
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker exec eiger-ollama-1 ollama run llama3.2:3b "Say hell
o"
.:
PS C:\Users\Harshu\OneDrive\Desktop\eiger> docker exec eiger-ollama-1 ollama run llama3.2:3b "Say hell
o"
Hello!

PS C:\Users\Harshu\OneDrive\Desktop\eiger> Invoke-WebRequest "http://127.0.0.1:8000/validate/m1?sessio
n=eiger-0e97c839" -UseBasicParsing | Select-Object -ExpandProperty Content
{"core":"pass","stretch":"fail"}
```

M2 After – Output Encoding Enabled

```
PS C:\Users\Harshu\OneDrive\Desktop\eiger> curl.exe "http://127.0.0.1:8000/validate/m2?session=eiger-0
e97c839"
{"core":"pass","stretch":"fail"}
PS C:\Users\Harshu\OneDrive\Desktop\eiger> $session="eiger-m2-fixed-" + (Get-Random)
PS C:\Users\Harshu\OneDrive\Desktop\eiger> $session
eiger-m2-fixed-635790952
PS C:\Users\Harshu\OneDrive\Desktop\eiger> $session="eiger-m2-fixed-635790952"
PS C:\Users\Harshu\OneDrive\Desktop\eiger> $payload = @{
>> session_id = $session
>> display_name = "<img src=x onerror=fetch('/beacon/xss?session=$session')>"
>> } | ConvertTo-Json
PS C:\Users\Harshu\OneDrive\Desktop\eiger>
PS C:\Users\Harshu\OneDrive\Desktop\eiger> Invoke-RestMethod -Uri "http://127.0.0.1:8000/api/profile"
-Method Post -ContentType "application/json" -Body $payload

status
-----
ok

PS C:\Users\Harshu\OneDrive\Desktop\eiger> curl.exe "http://127.0.0.1:8000/validate/m2?session=eiger-m
2-fixed-635790952"
{"core":"fail","stretch":"fail"}
PS C:\Users\Harshu\OneDrive\Desktop\eiger>
```

5. Validation Summary

Exercise	Before Fix	After Fix
----------	------------	-----------

M1 – Prompt Injection	Core: PASS	Core: FAIL
-----------------------	------------	------------

M2 – Output Handling/XSS	Core: PASS	Core: FAIL
--------------------------	------------	------------

The validation results demonstrate that the attacks were reproducible in vulnerable mode and that the corresponding security controls prevented the same attacks after remediation.

6. AI Usage

ChatGPT was used as a learning and analysis assistant during the assessment.

AI assistance was used for:

- Understanding the Eiger architecture and exercises
- Understanding prompt injection and output-encoding vulnerabilities
- Reviewing the relevant source code
- Developing safe test approaches for the authorized lab
- Troubleshooting Docker, WSL, and local environment issues
- Interpreting validation results
- Explaining security impact and mitigations
- Organizing screenshots and evidence
- Structuring the assessment report

All commands and security testing actions were reviewed and executed manually in the authorized Eiger local environment.

7. Conclusion

The Eiger assessment provided hands-on experience with two important AI/application security risks: **prompt injection** and **unsafe output handling**.

M1 demonstrated how attacker-controlled instructions can cause disclosure of sensitive information when secrets are exposed to the model.

M2 demonstrated how insufficient output encoding can allow stored XSS when user-controlled data is rendered as HTML.

Both vulnerabilities were reproduced in the vulnerable environment and successfully re-tested after remediation using the Eiger validation mechanism.

The assessment reinforced the importance of **separating trust boundaries, protecting secrets, encoding untrusted output, and validating security controls through adversarial testing**.